

# MongoDB

# Table of contents


---

1. Introduction
2. Query a MongoDB
  - a. Simple extraction and various simple commands
  - b. Aggregation
3. Data Modeling

# Chapter 1

# SQL and NoSQL evolution

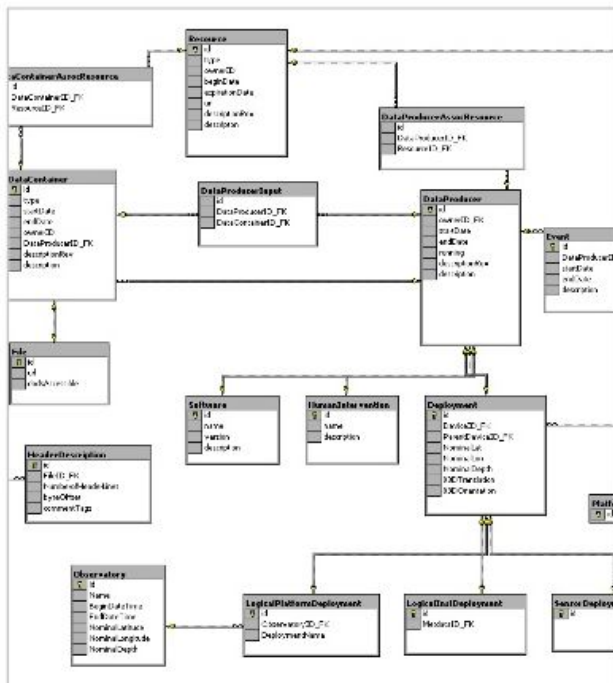
- 1970 : Dr. EF Codd published the paper “A Relation Model of Data for Large Shared Data Banks”
- 1974 : IBM developed SQL
- 1979 : Oracle provides the first RDBMS commercially available
- 1986 : The first SQL Standard is published
- 1998 : First time the term NoSQL is coined by Carlo Strozzi for an open source database that relaxes some SQL constraints (ACID) in favor of availability and scalability



# MongoDB key characteristics and use cases

Document-oriented database

RDBMS



MongoDB

```
{
  _id : ObjectId("4c4ba5e5e8aabbf3"),
  employee_name: "Dunham, Justin",
  department : "Marketing",
  title : "Product Manager, Web",
  report_up: "Neray, Graham",
  pay_band: "C",
  benefits : [
    { type : "Health",
      plan : "PPO Plus" },
    { type : "Dental",
      plan : "Standard" }
  ]
}
```

# MongoDB key characteristics and use cases

## Key characteristics

---

- General purpose database
  - Not the case for many other NoSQL databases (e.g. graph-oriented)
- Flexible schema design
  - Document oriented approaches with non-defined attributes
- Build for high availability
- Build to scale
- Aggregation framework
  - Provides powerful transformation capabilities
- Native replication
- Security features
- JSON based (BSON)
- Support MapReduce

# MongoDB key characteristics and use cases

What is the use case for MongoDB?

---

- Integration of data providing a single view of them
- Internet of Things
- Mobile applications
- Real-time analytics
- Personalization
- Catalog management
- Content management

# MongoDB key characteristics and use cases

## MongoDB criticism

---

- Schema-less nature
- Lack of proper ACID guarantees

⇒ Hard to ensure consistency, isolation...



# MongoDB key characteristics and use cases

## MongoDB binaries

---

- **mongod**
  - starts MongoDB
- **mongo**
  - open the MongoDB shell
- **mongodump**
  - create a database dump (bson+json files)
- **mongorestore**
  - restore a database dump

# Chapter 2

# CRUD using the shell

---

- List the databases
  - `db`
- Connect to a database (with creation if does not exist)
  - `use <database_name>`
- Insert a document
  - `db.<collection_name>.insert({...MATCHING_OBJECT...})`
  - `db.books.insert({title: 'mastering mongoDB', isbn: '101'})`
  - `db.books.insert({title: 'mongoDB cookbook', isbn: '1331'})`
- Find documents
  - `db.<collection_name>.find({...MATCHING_OBJECT...})`
  - `db.books.find({isbn: '1331'})`
- Delete a document
  - `db.<collection_name>.remove({...MATCHING_OBJECT...})`
  - `db.books.remove({isbn: '101'})`
- Delete a document (2)
  - `db.<collection_name>.remove(OBJECT_ID_STRING)`
  - `db.books.remove('2345678dae345eff')`

# CRUD using the shell

- Update a document
  - `db.<collection_name>.update({...MATCHING_OBJECT...},{...UPDATING_OBJECT...})`
  - `db.books.update({isbn: '1331'}, {price: 30})`  
**=> Replace all matched objects with the new object**
- Update a document (2)
  - `db.<collection_name>.update({...MATCHING_OBJECT...},{ $set: {...UPDATING_OBJECT...}})`
  - `db.books.update({isbn: '1331'}, { $set: {price: 30, title: 'Mongo'}})`
- Find all documents
  - `db.<collection_name>.find()`
- Find documents and pretty print them in the console
  - `db.<collection_name>.find({...MATCHING_OBJECT...}).pretty()`
  - `db.books.find().pretty()`

# Scripting for the mongo shell

---

Main reason for using the shell : Mongo shell is a Javascript shell

```
> var title = 'MongoDB in a nutshell'
> title
MongoDB in a nutshell
> db.books.insert({title: title, isbn: 102})
WriteResult({ "nInserted" : 1 })
> db.books.find()
{ "_id" : ObjectId("59203874141daf984112d080"), "title" : "MongoDB in a
nutshell", "isbn" : 102 }
```

# Scripting for the mongo shell

---

## Using scripts and functions

```
> queryBooksByIsbn = function(isbn) { return db.books.find({isbn: isbn})}  
> queryBooksByIsbn("101")  
{ "_id" : ObjectId("592035f6141daf984112d07f"), "title" : "mastering  
mongoDB", "isbn" : "101", "price" : 30 }
```

Executing script files:

From the shell

```
mongo FILE_PATH
```

From the mongo shell:

```
load(FILE_PATH_STRING)
```

# Scripting for the mongo shell

## Batch inserts using the shell

---

```
authorMongoFactory = function() {  
    for(loop=0;loop<1000;loop++){  
        db.books.insert({name: "MongoDB factory book" + loop})  
    }  
}
```

Although working, it performs 1000 database inserts (which is not efficient).

Prefer a bulk write:

```
fastAuthorMongoFactory = function() {  
    var bulk = db.books.initializeUnorderedBulkOp();  
    for(loop=0;loop<1000;loop++) {  
        bulk.insert({name: "MongoDB factory book" + loop})  
    }  
    bulk.execute();  
}
```

If you want the data having the same order of insertion than declaration, use

```
db.books.initializeOrderedBulkOp();
```

# Scripting for the mongo shell

Batch inserts using the shell

---

## Alternative

```
db.collection.bulkWrite(  
  [ <operation 1>, <operation 2>, ... ],  
    {  
      writeConcern : <document>,  
      ordered : <boolean>  
    }  
)
```



# Administration

---

- `db.dropDatabase()`
- `db.getCollectionNames()`
- `db.copyDatabase(fromDB, toDB)`

# Simple Extraction

# Simple extraction

find()

---

```
db.<collection>.find(  
    {...SELECTOR_OBJECT...},  
    {...PROJECTION_OBJECT...}  
)
```

# Simple extraction

`$exists`

---

`$exists`

allows to filter on properties existence

```
db.<collection>.find({<FIELD> : {$exists: <BOOL>}}, {})
```

# Simple extraction

## Projection

---

The projection allows to get specific property from the objects

```
db.<collection>.find({}, {<FIELD> : <BOOL>})
```

When <BOOL> is

- true, then shows only the specified properties
- false, then hides the specified properties

# Simple extraction

\$eq, \$gt, \$lt, \$gte, \$lte, \$ne

Comparison operators

```
db.<collection>.find({  
  <FIELD>: {<OPERATOR>:<VALUE>}  
}, {})
```

*Example:*

```
db.products.find({price: {$lte: 50}}, {})
```

\$eq	equal
\$gt	greater than
\$lt	lower than
\$gte	greater equal
\$lte	lower equal
\$ne	not equal

# Simple extraction

\$in, \$nin

---

**Is in** or **is not in** operator

```
db.<collection>.find({  
    <FIELD>: {<OPERATOR>:<VALUES_ARRAY>}  
}, {})
```

*Example:*

```
db.customers.find({City: {$in: ['Orlando', 'Baton Rouge']}})
```

# Simple extraction

\$and, \$or, \$nor, \$not

## Logical operators

```
db.<collection>.find({  
  <OPERATOR>: [<EXPR1>, <EXPR2>, ...]  
}, {})
```

### Example:

```
db.products.find({$and: [  
  {City: 'Orlando'}, {Price: {$lt:50}}  
]}, {})
```

\$and	Joins query clauses with a logical AND returns all documents that match the conditions of both clauses.
\$or	Joins query clauses with a logical OR returns all documents that match the conditions of either clause.
\$nor	Joins query clauses with a logical NOR returns all documents that fail to match both clauses.
\$not	



# Simple extraction

\$type

---

Selects the documents where the value of the field is an instance of the specified BSON type(s).

```
db.<collection>.find({  
  <FIELD>: { $type : <YPEESTRING> }  
}, {})
```

## *Example:*

```
db.products.find({ $and: [  
  { Price: { $type: 'double' } }  
] }, {})
```

# Simple extraction

\$mod

---

Select documents where the value of a field divided by a divisor has the specified remainder

```
db.<collection>.find({  
  <FIELD>: { $mod : [<DIVISOR>, <REMAINDER>] }  
}, {})
```

## *Example:*

```
db.products.find({ $and: [  
  { Price: { $mod: [2, 0] } }  
] }, {})
```

# Simple extraction

\$regex

Provides regular expression capabilities for pattern matching strings in queries. (PCRE 8.41)

```
db.<collection>.find({  
  <FIELD>: {$regex : /REGEX/, $options : <STRING_OPTIONS>}  
}, {})
```

```
db.<collection>.find({  
  <FIELD>: /REGEX/OPTIONS  
}, {})
```

*Example:*

```
db.products.find({productName: /cup/i})
```

i	Case insensitivity to match upper and lower cases.
m	For patterns that include anchors (i.e. ^ for the start, \$ for the end), match at the beginning or end of each line for strings with multiline values. Without this option, these anchors match at beginning or end of the string.
x	“Extended” capability to ignore all white space characters in the \$regex pattern unless escaped or included in a character class.

# Simple extraction

\$text

---

Performs a text search on the content of the fields indexed with a text index.

```
REQUIRED: db.<collection>.createIndex ({<FIELD>:"text"})
db.<collection>.find({
  $text: {
    $search: <TEXT_TO_SEARCH>
    $language: <LANGUAGE>
    $caseSensitive: <BOOL>
    $diacriticSensitive: <BOOL>
  }
}, {}))
```

# Simple extraction

\$text

---

Performs a text search on the content of the fields indexed with a text index.

## *Example:*

```
db.products.find({
  $text: {
    $search: 'fry',
    $language: 'english',
    $diacriticSensitive: true
  }
})
```

# Simple extraction

\$where

---

Use the [\\$where](#) operator to pass either a string containing a JavaScript expression or a full JavaScript function to the query system. (From 3.6, \$expr should be preferred).

```
db.products.find({  
  $where: <JS_EXPRESSION>  
})
```

*Example:*

```
db.embeddedOrders.find({$where: 'this.Orders.length == 1'})
```

# Simple extraction

\$expr (3.6)

Allows the use of aggregation expressions within the query language. [\\$expr](#) can build query expressions that compare fields from the same document in a [\\$match](#) stage

```
db.products.find({  
  $expr: <EXPRESSION>  
})
```

## *Example:*

```
db.monthlyBudget.find({$expr:{ $gt: ["$spent", "$budget"] }})
```

# Simple extraction

`.distinct()`

---

Finds the distinct values for a specified field across a single collection or view and returns the results in an array.

```
db.collection.distinct(  
    field,  
    query  
)
```



# Simple extraction

`.count()`

---

Returns the count of documents that would match a `find()` query for the collection or view

```
db.collection.count(  
    query,  
    options  
)
```

# Simple extraction

## Exercices

---

1. Get the customers without company
2. List distinctly all the city where customers working for “Yodo” or “Kare” are living
3. Count the number of products which are sold at a price greater than 50
4. Display the first and last name of the customer having a \*\*\*\*\*@patch.com email adres.
5. Find purple product costing more than 20
6. Find the product having their color in their name

# Aggregation

# Aggregation

\$expr (3.6)

---

```
db.<collection>.aggregate ([  
    {<EXPRESSION1>},  
    {<EXPRESSION2>},  
    {<EXPRESSION3>},  
    ...  
])
```

Aggregation is performed step-by-step. The document resulting from an expression is used as input for the following expression.

# Aggregation Pipeline Stages

\$match

---

Filters the documents to pass only the documents that match the specified condition(s) to the next pipeline stage.

```
db.<collection>.aggregate([  
    {$match: <QUERY>},  
    ...  
])
```

*Example:*

```
db.products.aggregate([{$match: {price: {$gt:50}}}]])
```

# Aggregation Pipeline Stages

## \$group

Groups documents by some specified expression and outputs to the next stage a document for each distinct grouping. The output documents contain an `_id` field which contains the distinct group by key.

```
db.<collection>.aggregate([
  {$group: {
    _id: <EXPRESSION>,
    <FIELD1>: { <ACC1> : <EXPRESSION1> },
    <FIELD2>: { <ACC2> : <EXPRESSION2> },
    ...
  }},
  ...
])
```

\$avg
\$first
\$last
\$max
\$min
\$push
\$addToSet
\$stdDevPop
\$stdDevSamp
\$sum

# Aggregation Pipeline Stages

\$group

Groups documents by some specified expression and outputs to the next stage a document for each distinct grouping. The output documents contain an `_id` field which contains the distinct group by key.

## *Example:*

```
db.<collection>.aggregate([
  {$group: {
    _id: <EXPRESSION>,
    <FIELD1>: { <ACC1> : <EXPRESSION1> },
    <FIELD2>: { <ACC2> : <EXPRESSION2> },
    ...
  }},
  ...
])
```

# Aggregation Pipeline Stages

## \$unwind

---

Deconstructs an array field from the input documents to output a document for each element. Each output document is the input document with the value of the array field replaced by the element.

```
db.<collection>.aggregate([
{
  $unwind:
    {
      path: <FIELD_PATH>,
      includeArrayIndex: <STRING>,
      preserveNullAndEmptyArrays: <BOOL>
    }
},
...
])
```

**To specify a field path, prefix the field name with a dollar sign \$ and enclose in quotes.**



# Aggregation Pipeline Stages

## \$lookup

---

Performs a left outer join to an unsharded collection in the *same* database to filter in documents from the “joined” collection for processing.

```
db.<collection>.aggregate([
{
  $lookup:
  {
    from: <collection to join>,
    localField: <field from the input documents>,
    foreignField: <field from the documents of the "from" collection>,
    as: <output array field>
  }
},
...
])
```

# Aggregation Pipeline Stages

## \$sort

---

Sorts all input documents and returns them to the pipeline in sorted order.

```
db.<collection>.aggregate([
{
    $sort: {
        <field1>: <sort order>,
        <field2>: <sort order>,
        ...
    }
},
...
])
```

# Aggregation Pipeline Stages

## \$limit

Limits the number of documents passed to the next stage in the [pipeline](#).

```
db.<collection>.aggregate([
{
    $limit: <POSITIVE_NUMBER>
},
...
])
```

### NOTE:

When a `$sort` precedes a `$limit` and there are no intervening stages that modify the number of documents, the optimizer can coalesce the `$limit` into the `$sort`. This allows the `$sort` operation to only maintain the top `n` results as it progresses, where `n` is the specified limit, and ensures that MongoDB only needs to store `n` items in memory. This optimization still applies when `allowDiskUse` is `true` and the `n` items exceed the [aggregation memory limit](#).

# Aggregation Pipeline Stages

\$out

---

Takes the documents returned by the aggregation pipeline and writes them to a specified collection. The \$out operator must be *the last stage* in the pipeline.

```
db.<collection>.aggregate([
{
    $out: {<OUTPUT_COLLECTION_STRING>}
},
    ...
])
```

# Aggregation Pipeline Stages

## Exercises

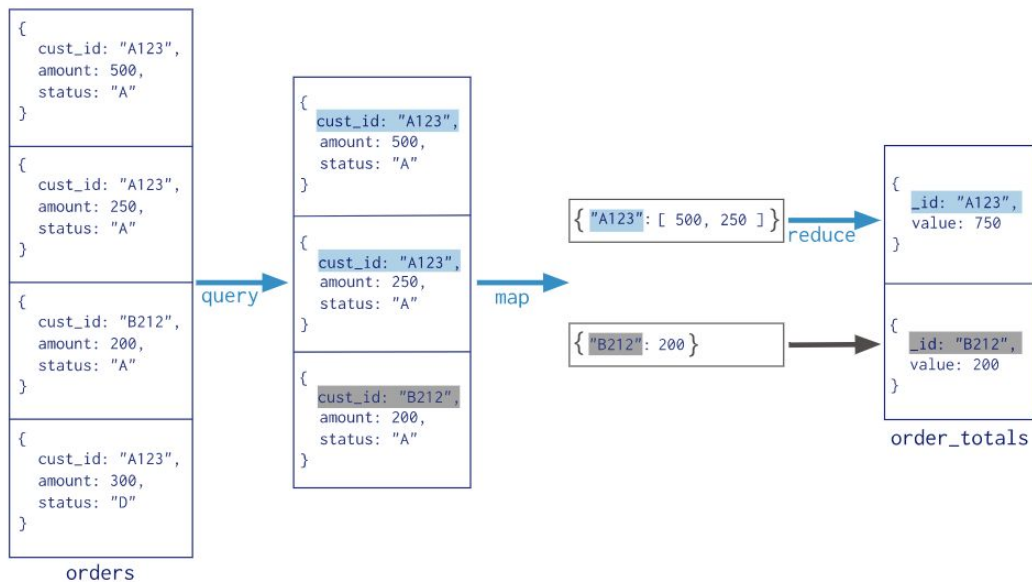
---

- Create a new collection which is denormalized where:*
  - Customers have Orders that are composed of OrderLine*
- Count the number of orders for each customer working at Yodo
  - Display only the first name, the last name and the orders number
- Display the sales amount and the quantity ordered for each product ordered by Yodo employees
- What is the product name of the product the most bought by women*
- What is the average order value for the people living in Atlanta
- Determine the sales amount both by color and by city*
- For each customer living in Orlando, determines how many orders included leaf related products.*
- Compare the sales amount of cherry and tomato related products*

# MapReduce

# MapReduce

Collection  
↓  
db.orders.mapReduce(  
  map    → function() { emit( this.cust\_id, this.amount ); },  
  reduce → function(key, values) { return Array.sum( values ); },  
  query  → { query: { status: "A" },  
  output → out: "order\_totals"  
  }  
)



# MapReduce

## Exercise

---

1 - Find the average number of product (orderlines) each customer of Yodo ordered per order

+ Exercices 3 and 6 from *Aggregation Pipeline Stages* exercises



# Chapter 3

# Data Modelling

## Flexible schema

---

Unlike SQL databases, MongoDB collections do not require the documents to have the same schema.

- Fields, data type... can vary from one document to another.

In practice, the documents of a particular collection should share a similar schema

- Those similarities can be ensured by rules. This is called schema validation.

# Data Modelling

## Document structure

---

- Relational databases are designed on the basis of the real-world relations
- Document-oriented database are designed on the basis of how that application will use those data

MongoDB allows both the use of embedded data (denormalisation) or references (normalisation).

# Data Modelling

## Embedded data

**Embedded data** is a kind a denormalisation.

**Purpose:** allowing the application to retrieve and manipulate related-data in a single operation.

### Use:

- One-to-one “contains” relationship between data  
Ex: ID card for a person...
- When a child document is always viewed in the context of the parent  
Ex: A customer list of delivery address

```
{
  _id: <ObjectId>,
  username: "123xyz",
  contact: {
    phone: "123-456-7890",
    email: "xyz@example.com"
  },
  access: {
    level: 5,
    group: "dev"
  }
}
```

Embedded sub-document

Embedded sub-document

# Data Modelling

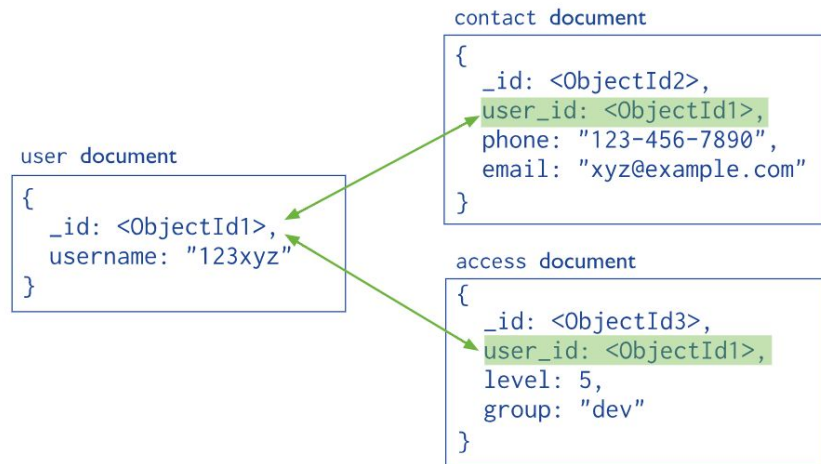
## References

**References** are used to normalise.

**Purpose:** provides more flexibility and avoid data duplication.

**Use:**

- when embedding would result in duplication of data
- to represent more complex many-to-many relationships
- to model large hierarchical data sets



# Data Modeling

## Rules of Thumb for MongoDB Schema Design

---

Modeling **One-to-Few** : Embedded documents

*Example: Person ↔ Address*

*Addresses are embedded in the person*

Modeling **One-to-Many** : References (Many ~ hundreds)

The parent is referencing the children

*Example: Product ↔ Parts*

*Product having an array with references to all parts*

Modeling **One-to-Squillions** : References (Many ~squillions)

Each child has a reference to its parent (otherwise the size of the document will explode)

*Example: Host ↔ Log*

*Each Log document has a reference to its host*

# Data Modeling

## JSON Schema

---

### Add validation at collection creation:

```
db.createCollection(<collection_name>, {  
  validator: {  
    $jsonSchema: {  
    }  
  }  
})
```

### Add validation when the collection exists:

```
db.runCommand( { collMod: <collection_name>, validator: {  
  $jsonSchema: {  
  
  }  
} })
```

# Library System

## Exercise

---

Assume there is a library system with the following properties.

The library contains one or several copies of the same book. Every copy of a book has a copy number and is located at a specific location in a shelf. A copy is identified by the copy number and the ISBN number of the book. Every book has a unique ISBN, a publication year, a title, an author, and a number of pages. Books are published by publishers. A publisher has a name as well as a location.

Within the library system, books are assigned to one or several categories. A category can be a subcategory of exactly one other category. A category has a name and no further properties. Each reader needs to provide his/her family name, his/her first name, his/her city, and his/her date of birth to register at the library. Each reader gets a unique reader number. Readers borrow copies of books. Upon borrowing the return date is stored.

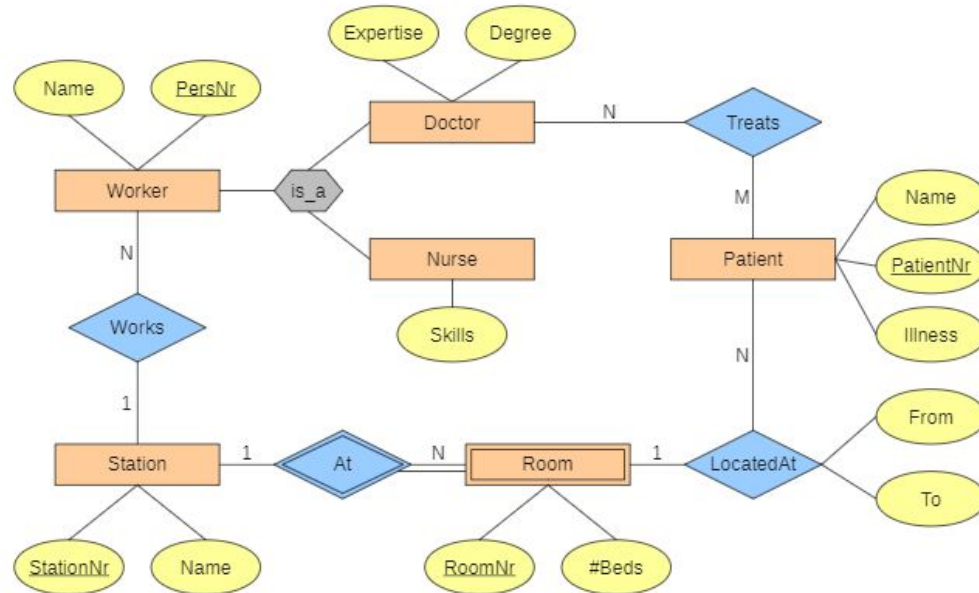
How would you model this situation in a document-oriented database ?



# Data Modeling

## Exercise 2

Transform this ER Model into document schemas



# Data Modeling

## Exercise 3

---

Implement the following document schema:

Customers are always known by their first name and last name that are strings. We can know their address which consists of a street, a zip code (int varying from 1000 to 9999) and a city.

There are no other possible field.